

AI Introduction — Tutorial

Use any common AI chat tool (ChatGPT, Claude, Gemini, Copilot, DeepSeek, or similar). You do not need to write code.

Task 1 — Explore interesting tasks of an AI agent

Time: 50 minutes

Look at AI agents doing real work in an environment: playing games, handling office tasks, and so on. The point is not which demo looks most impressive. Notice what the agent can observe, what it can actually do, and where a fluent plan hides a failed action.

Example video: [AI Simulates Civilization in Minecraft](#)

Three large language models (GPT, Claude, Gemini) are given Minecraft bodies. They are not only *talking* about the game. They chop wood, dig, build, try Redstone, and face a zombie apocalypse. Who wins matters less than this:

Once an AI has a body, tools, and lasting consequences, where does it look competent, and where does a good explanation hide a failed action?

You may also look at other agent demos (office work, browsing, coding, games). For each example, jot down:

1. What is the goal?
2. What can the agent see or use?
3. One moment that looked smart, and one moment that failed.

Task 2 — Ask the agent what it can and cannot do

Time: 50 minutes

Use the **same model** and the **same identity** for every prompt. Do not stop after one vague question.

Step 0 — Write your identity

Complete this before you start. Keep it unchanged for all prompts.

I am _____ (major / role / club position).
The one concrete task I have this week is _____.
If this work goes wrong, it would harm _____.

Do not paste: ID numbers, unpublished grades, confidential company data, client lists, passwords, medical records, or other people's private information.

Run at least four prompts

Save the full prompt and the full reply (or a screenshot).

Prompt A — empty baseline

What can you do for me? What can you not do?

Prompt B — identity, still no task

I am a [your role]. I have limited time this week.
What kinds of work can you usefully do for me?
What should I never outsource to you?
Please be concrete. If you are not sure, say you are not sure.

Prompt C — one real task, lazy instruction

Help me with [the one task you wrote above].

Watch for invented dates, names, rules, or policies.

Prompt D — the same task, executable instruction

Task: [one sentence]
Audience: [who will use the result]
Known facts (do not invent beyond this list):
- ...
- ...
Constraints:
- Do not invent names, dates, rules, or numbers I did not give.
- If something is missing, label it UNKNOWN and ask me.
- Length / format: [e.g. 8 bullets / a 150-word email]
After the draft, list:
1) what you assumed
2) what I must check before I use this

Prompt E — boundary probes (choose two)

1. Ask for a school rule you did **not** provide (for example, who must sign a three-day leave).
2. Ask it to book a room on Friday and send you the confirmation number.
3. Give only “I have homework next week” and see whether it invents a full calendar.
4. Ask it to mark its own draft and *prove* every fact is correct.

A refusal can be a good result. A fluent invention is a bad result.

Comparison table (required)

Round	What I changed in the prompt	Clearest change in the output	Something it claimed that I do not trust	More useful / more dangerous / no difference
A				
B				
C				
D				
E				